

```
#calling packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import pyfixest as pf
#importing data
tax_data = pd.read_csv('/Users/avanthpakanati/Desktop/ECON:HLTH Research Seminar /Homework3/')
```

ModuleNotFoundError: No module named 'pyfixest'

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
Input In [6], in <cell line: 7>()
      5 import seaborn as sns
      6 import statsmodels.api as sm
----> 7 import pyfixest as pf
      8 #importing data
      9 tax_data = pd.read_csv('/Users/avanthpakanati/Desktop/ECON:HLTH Research Seminar /Homework3/')
```

ModuleNotFoundError: No module named 'pyfixest'

```
#summarize the data
#question 1


#finding prop of states that had change in tax from 1970-1986
tax_data = tax_data.sort_values(by=['state', 'Year'])
tax_data_1 = tax_data[(tax_data['Year']>=1970) & (tax_data['Year']<=1985)]
tax_data_1['tax_change'] = tax_data_1.groupby('state')['tax_state'].diff().ne(0).astype(int)
tax_change = tax_data_1.groupby('Year')['tax_change'].mean()

plt.bar(tax_change.index, tax_change.values)
plt.title('Proportion of States with Change in Cigarette Tax (1970-1985)')
plt.xlabel('Year')
plt.show()
```

/var/folders/2q/wzjp_2kd355b8clhzqwmtytb40000gn/T/ipykernel_2702/86190851.py:7: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide

```
tax_data_1['tax_change'] = tax_data_1.groupby('state')['tax_state'].diff().ne(0).astype(int)
```




```
#question 2
#in 2012 $$$
cpi_2012 = tax_data.loc[tax_data['Year'] == 2012, 'price_cpi'].iloc[0]
tax_data['tax_dollar_2012'] = tax_data['tax_dollar'] * (cpi_2012 / tax_data['price_cpi'])
tax_data['price_per_pack_2012'] = tax_data['cost_per_pack'] * (cpi_2012 / tax_data['price_cpi'])

avg_values = tax_data.groupby('Year')[['tax_dollar_2012', 'price_per_pack_2012']].mean()

#plot graph
plt.plot(avg_values.index, avg_values['tax_dollar_2012'], label='Average Tax (2012 dollars)')
plt.plot(avg_values.index, avg_values['price_per_pack_2012'], label='Average Price (2012 dollars)')
plt.legend()
plt.title('Avg Cigarette Tax and Price (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Dollars (2012)')
```

```
plt.grid(True)
plt.show()
```



```
avg_values = tax_data.groupby('Year')[['tax_2012', 'price_cpi']].mean()

#line graph
plt.plot(avg_values.index, avg_values['tax_2012'], label='Average Tax (2012 dollars)')
plt.plot(avg_values.index, avg_values['price_cpi'], label='Average Price (2012 dollars)')
plt.legend()
plt.title('Cigarette Tax and Price in 2012 Dollars (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Dollars (2012 value)')
plt.grid(True)
plt.show()
```

KeyError: "Columns not found: 'tax_2012'"

```

-----
KeyError                                Traceback (most recent call last)
Input In [5], in <cell line: 1>()
----> 1 avg_values = tax_data.groupby('Year')[['tax_2012', 'price_cpi']].mean()
      4 #line graph
      5 plt.plot(avg_values.index, avg_values['tax_2012'], label='Average Tax (2012 dollars)')
File ~/anaconda/lib/python3.9/site-packages/pandas/core/groupby/generic.py:1338, in DataFrameGroupBy.getitem
    1329 if isinstance(key, tuple) and len(key) > 1:
    1330     # if len == 1, then it becomes a SeriesGroupBy and this is actually
    1331     # valid syntax, so don't raise warning
    1332     warnings.warn(
    1333         "Indexing with multiple keys (implicitly converted to a tuple of
    1334         keys) will be deprecated, use a list instead.",
    1335         FutureWarning,
    1336         stacklevel=find_stack_level(),
    1337     )
-> 1338 return super().__getitem__(key)
File ~/anaconda/lib/python3.9/site-packages/pandas/core/base.py:240, in SelectionMixin.__getitem__
    238     if len(self.obj.columns.intersection(key)) != len(set(key)):
    239         bad_keys = list(set(key).difference(self.obj.columns))
--> 240         raise KeyError(f"Columns not found: {str(bad_keys)[1:-1]}")
    241     return self._getitem(list(key), ndim=2)
    243 elif not getattr(self, "as_index", False):
KeyError: "Columns not found: 'tax_2012'"

```

```

#question 3
#years 1970-2018
tax_data_3 = tax_data[(tax_data['Year'] >= 1970) & (tax_data['Year'] <= 2018)]

change_in_price = tax_data_3.groupby('state')['price_cpi'].agg(['first', 'last'])
change_in_price['price_increase'] = change_in_price['last'] - change_in_price['first']


#finding 5 states with highest price increase
top_5_states = change_in_price.nlargest(5, 'price_increase').index.tolist()
top_5 = tax_data_3[tax_data_3['state'].isin(top_5_states)]

#plotting the graph

plt.figure(figsize=(10, 6))
for state in top_5_states:
    state_data = top_5[top_5['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)

```

```
plt.title('Average Packs Sold Per Capita in States with 5 Highest Price Increases (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Packs Sold Per Capita')
plt.legend()
plt.grid(True)
plt.show()
```




```
#question 4
#Lowest increase in cig prices

# 5 states with the lowest price increase
bottom_5_states = change_in_price.nsmallest(5, 'price_increase').index.tolist()
bottom_5 = tax_data_3[tax_data_3['state'].isin(bottom_5_states)]

plt.figure(figsize=(10, 6))
for state in bottom_5_states:
    state_data = bottom_5[bottom_5['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)
```


```
plt.title('Average Packs Sold Per Capita in States with 5 Lowest Price Increases (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Packs Sold Per Capita')
plt.legend()
plt.grid(True)
plt.show()
```



```
#question 5 comparing states with lowest and highest price increase
#putting it all in one graph
plt.figure(figsize=(9, 7))
for state in top_5_states:
    state_data = top_5[top_5['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)
for state in bottom_5_states:
    state_data = bottom_5[bottom_5['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)

plt.title('Avg Packs Sold Per Capita in States With Highest/Lowest Price Increases (1970-2018)')
```

```
plt.xlabel('Year')
plt.ylabel('Packs Sold Per Capita')
plt.legend()
plt.grid(True)
plt.show()
```



```
#Estimating ATE
```

```
#QUESTION 6
```

```
cig_data = tax_data[(tax_data['Year'] >= 1970) & (tax_data['Year'] <= 1990)]
cig_data['ln_sales'] = np.log(cig_data['sales_per_capita'])
cig_data['ln_total_tax'] = np.log(cig_data['tax_dollar'])
cig_data['ln_price'] = np.log(cig_data['price_cpi'])
```

```
# OLS
X = sm.add_constant(cig_data['ln_price']) # Add constant for intercept
Y = cig_data['ln_sales']

model = sm.OLS(Y, X).fit()
print(model.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  ln_sales      R-squared:                0.294
Model:                            OLS      Adj. R-squared:            0.293
Method:                 Least Squares      F-statistic:                445.1
Date:                Tue, 18 Mar 2025      Prob (F-statistic):        6.98e-83
Time:                  13:29:17      Log-Likelihood:            263.40
No. Observations:                1071      AIC:                       -522.8
Df Residuals:                    1069      BIC:                       -512.8
Df Model:                            1
Covariance Type:                nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const           5.3854      0.028    193.692      0.000        5.331        5.440
ln_price       -0.8094      0.038   -21.098      0.000       -0.885       -0.734
=====
Omnibus:                 89.160   Durbin-Watson:           0.183
Prob(Omnibus):            0.000   Jarque-Bera (JB):        466.536
Skew:                    0.128   Prob(JB):                4.93e-102
Kurtosis:                 6.223   Cond. No.                 10.0
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
/var/folders/2q/wzjp_2kd355b8clhzqwmytb40000gn/T/ipykernel_2702/707357993.py:2: SettingWithC
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide
cig_data['ln_sales'] = np.log(cig_data['sales_per_capita'])
```



```
/var/folders/2q/wzjp_2kd355b8clhzqwmtytb40000gn/T/ipykernel_2702/707357993.py:3: SettingWithC
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guid
cig_data['ln_total_tax'] = np.log(cig_data['tax_dollar'])
/var/folders/2q/wzjp_2kd355b8clhzqwmtytb40000gn/T/ipykernel_2702/707357993.py:4: SettingWithC
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guid
cig_data['ln_price'] = np.log(cig_data['price_cpi'])
```

```
#Question 7 and 8, log sales on log prices
#First stage
first_stage = sm.OLS(cig_data['ln_price'], sm.add_constant(cig_data['ln_total_tax'])).fit()
print("First stage Regression (ln_price ~ ln_total_tax):\n")
print(first_stage.summary())
```

First stage Regression (ln_price ~ ln_total_tax):

```

                                OLS Regression Results
=====
Dep. Variable:                  ln_price    R-squared:                 0.683
Model:                            OLS      Adj. R-squared:            0.683
Method:                 Least Squares    F-statistic:                 2301.
Date:                   Mon, 17 Mar 2025    Prob (F-statistic):          8.21e-269
Time:                   09:08:01      Log-Likelihood:             -86.164
No. Observations:          1071      AIC:                        176.3
Df Residuals:              1069      BIC:                        186.3
Df Model:                   1
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1.1786	0.033	35.712	0.000	1.114	1.243
ln_total_tax	1.0803	0.023	47.973	0.000	1.036	1.125

```

=====
Omnibus:                 30.760    Durbin-Watson:                 0.408
Prob(Omnibus):            0.000    Jarque-Bera (JB):             32.668
Skew:                     0.421    Prob(JB):                     8.06e-08
Kurtosis:                 3.156    Cond. No.                     8.72
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
# Log prices from the first stage
price_hat = first_stage.predict(sm.add_constant(cig_data['ln_total_tax']))

# Second-stage regression (IV Regression)
second_stage = sm.OLS(cig_data['ln_sales'], sm.add_constant(price_hat)).fit()
print("\nSecond stage Regression (ln_sales ~ pricehat):\n")
print(second_stage.summary())
```

Second stage Regression (ln_sales ~ pricehat):

```

                        OLS Regression Results
=====
Dep. Variable:          ln_sales      R-squared:          0.236
Model:                  OLS          Adj. R-squared:       0.235
Method:                 Least Squares  F-statistic:        330.3
Date:                  Mon, 17 Mar 2025  Prob (F-statistic):  1.56e-64
Time:                  09:08:01      Log-Likelihood:     221.17
No. Observations:      1071         AIC:                -438.3
Df Residuals:          1069         BIC:                -428.4
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	4.7101	0.008	573.443	0.000	4.694	4.726
0	-0.2843	0.016	-18.175	0.000	-0.315	-0.254

```
=====
Omnibus:                83.338      Durbin-Watson:        0.157
Prob(Omnibus):          0.000      Jarque-Bera (JB):     430.014
Skew:                   0.023      Prob(JB):             4.20e-94
Kurtosis:               6.104      Cond. No.             2.98
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
#Question 9 - repeat for 1991-2015
cig_data2 = tax_data[(tax_data['Year'] >= 1991) & (tax_data['Year'] <= 2015)]

cig_data2['ln_sales'] = np.log(cig_data2['sales_per_capita'])
cig_data2['ln_price'] = np.log(cig_data2['cost_per_pack'])
cig_data2['ln_total_tax'] = np.log(cig_data2['tax_dollar'])

#running OLS regression
X2 = sm.add_constant(cig_data2['ln_price'])
Y2 = cig_data2['ln_sales']

reg2 = sm.OLS(Y2, X2).fit()
print(reg2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  ln_sales      R-squared:                0.533
Model:                            OLS      Adj. R-squared:           0.532
Method:                 Least Squares      F-statistic:              1451.
Date:                Mon, 17 Mar 2025      Prob (F-statistic):       1.52e-212
Time:                09:08:01      Log-Likelihood:          -296.47
No. Observations:          1275      AIC:                     596.9
Df Residuals:              1273      BIC:                     607.2
Df Model:                    1
Covariance Type:            nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const           5.0395      0.023     219.934      0.000        4.995        5.084
ln_price        -0.6656      0.017    -38.094      0.000       -0.700       -0.631
=====
Omnibus:                 19.351   Durbin-Watson:           0.158
Prob(Omnibus):            0.000   Jarque-Bera (JB):        33.046
Skew:                     0.064   Prob(JB):                6.67e-08
Kurtosis:                 3.778   Cond. No.                 5.37
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

/var/folders/2q/wzjp_2kd355b8clhzqmytb40000gn/T/ipykernel_90894/1083457502.py:4: SettingWith

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide
`cig_data2['ln_sales'] = np.log(cig_data2['sales_per_capita'])`
/var/folders/2q/wzjp_2kd355b8clhzqwmtytb40000gn/T/ipykernel_90894/1083457502.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide
`cig_data2['ln_price'] = np.log(cig_data2['cost_per_pack'])`
/var/folders/2q/wzjp_2kd355b8clhzqwmtytb40000gn/T/ipykernel_90894/1083457502.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide
`cig_data2['ln_total_tax'] = np.log(cig_data2['tax_dollar'])`

```
#QUESTION 10
# Creating a table to summarize results
summary_table = pd.DataFrame({
    'Years': ['1970-1990', '1970-1990', '1991-2015', '1991-2015'],
```

ValueError: All arrays must be of the same length

```
-----
ValueError                                Traceback (most recent call last)
Input In [29], in <cell line: 3>()
      1 #QUESTION 10
      2 # Creating a table to summarize results
----> 3 results_table = pd.DataFrame(          4      'Time Period': ['1970-1990', '1970-1990', '1991-2015', '1991-2015'],
      13 # Print the table without the index and with lines between each column/row
      14 print(results_table.to_string(index=False, line_width=80))
File ~/anaconda/lib/python3.9/site-packages/pandas/core/frame.py:636, in DataFrame.__init__(self, data, index, columns, dtype, copy)
    630     mgr = self._init_mgr(
    631         data, axes="index": index, "columns": columns, dtype=dtype, copy=copy
    632     )
    634 elif isinstance(data, dict):
    635     # GH#38939 de facto copy defaults to False only in non-dict cases
--> 636     mgr = dict_to_mgr(data, index, columns, dtype=dtype, copy=copy, typ=manager)
    637 elif isinstance(data, ma.MaskedArray):
    638     import numpy.ma.mrecords as mrecords
File ~/anaconda/lib/python3.9/site-packages/pandas/core/internals/construction.py:502, in dict_to_mgr(data, index, columns, dtype, copy)
```

```

494     arrays = [
495         x
496         if not hasattr(x, "dtype") or not isinstance(x.dtype, ExtensionDtype)
497         else x.copy()
498         for x in arrays
499     ]
500     # TODO: can we get rid of the dt64tz special case above?
--> 502 return arrays_to_mgr(arrays, columns, index, dtype=dtype, typ=typ, consolidate=copy)
File ~/anaconda/lib/python3.9/site-packages/pandas/core/internals/construction.py:120, in ar
117 if verify_integrity:
118     # figure out the index, if necessary
119     if index is None:
--> 120         index = _extract_index(arrays)
121     else:
122         index = ensure_index(index)
File ~/anaconda/lib/python3.9/site-packages/pandas/core/internals/construction.py:674, in _e
672 lengths = list(set(raw_lengths))
673 if len(lengths) > 1:
--> 674     raise ValueError("All arrays must be of the same length")
676 if have_dicts:
677     raise ValueError(
678         "Mixing dicts with non-Series may lead to ambiguous ordering."
679     )
ValueError: All arrays must be of the same length

```