

```

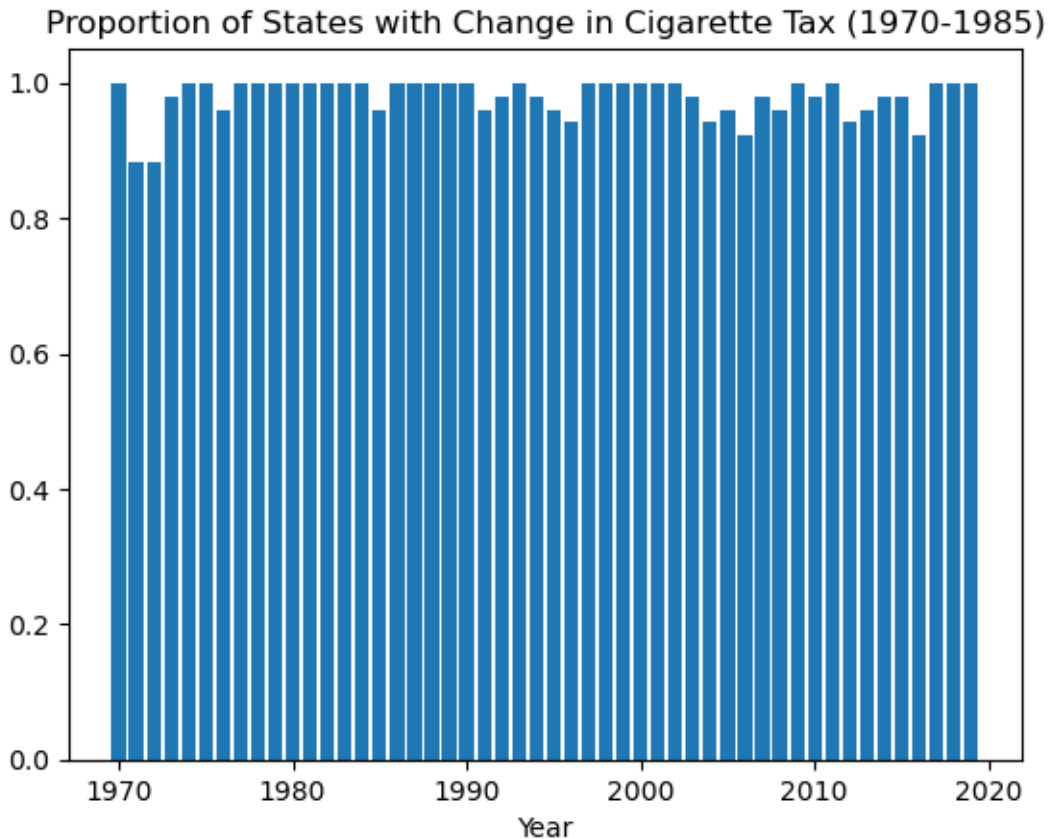
#calling packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

#importing data
tax_data = pd.read_csv('/Users/avanthpakanati/Desktop/ECON:HLTH Research Seminar /Homework3/

#summarize the data
#question 1
tax_data = tax_data.sort_values(by=['state', 'Year'])
tax_data['tax_change'] = tax_data.groupby('state')['tax_percent'].diff().ne(0).astype(int)
tax_change_proportion = tax_data.groupby('Year')['tax_change'].mean()

plt.bar(tax_change_proportion.index, tax_change_proportion.values)
plt.title('Proportion of States with Change in Cigarette Tax (1970-1985)')
plt.xlabel('Year')
plt.show()


```



```
#question 2
#in 2012 $$$
cpi_2012 = tax_data.loc[tax_data['Year'] == 2012, 'price_cpi'].iloc[0]
tax_data['tax_dollar_2012'] = tax_data['tax_dollar'] * (cpi_2012 / tax_data['price_cpi'])
tax_data['price_per_pack_2012'] = tax_data['cost_per_pack'] * (cpi_2012 / tax_data['price_cpi'])

avg_values = tax_data.groupby('Year')[['tax_dollar_2012', 'price_per_pack_2012']].mean()

#plot graph
plt.plot(avg_values.index, avg_values['tax_dollar_2012'], label='Average Tax (2012 dollars)')
plt.plot(avg_values.index, avg_values['price_per_pack_2012'], label='Average Price (2012 dollars)')
plt.legend()
plt.title('Avg Cigarette Tax and Price (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Dollars (2012 value)')
plt.grid(True)
plt.show()
```



```
#Question 3
# identify 5 states w/ highest inncrease in cig prices
tax_data_2018 = tax_data[tax_data['Year'] == 2018].set_index('state')
tax_data_1970 = tax_data[tax_data['Year'] == 1970].set_index('state')


price_increase_state = tax_data_2018['cost_per_pack'] - tax_data_1970['cost_per_pack']
top_5_states = price_increase_state.nlargest(5).index
top_5_data = tax_data[tax_data['state'].isin(top_5_states)]

plt.figure(figsize=(9, 7))
for state in top_5_states:
    state_data = top_5_data[top_5_data['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)

plt.title('Average Packs Sold Per Capita in States With Highest Price Increases (1970-2018)')
plt.xlabel('Year')
```

```
plt.ylabel('Packs Sold Per Capita')
plt.legend()
plt.grid(True)

plt.show()
```



```
#question 4
#Lowest increase in cig prices


# 5 states with the lowest price increase
bottom_5_states = price_increase_state.nsmallest(5).index
bottom_5_data = tax_data[tax_data['state'].isin(bottom_5_states)]
```

```

plt.figure(figsize=(9, 7))
for state in bottom_5_states:
    state_data = bottom_5_data[bottom_5_data['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)

plt.title('Average Packs Sold Per Capita in States with Lowest Price Increases (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Packs Sold Per Capita')
plt.legend()
plt.grid(True)
plt.show()

```



```

#question 5 comparing states with lowest and highest price increase
plt.figure(figsize=(9, 7))


```

```

for state in top_5_states:
    state_data = top_5_data[top_5_data['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)
for state in bottom_5_states:
    state_data = bottom_5_data[bottom_5_data['state'] == state]
    plt.plot(state_data['Year'], state_data['sales_per_capita'].rolling(window=3).mean(), label=state)

plt.title('Avg Packs Sold Per Capita in States With Highest/Lowest Price Increases (1970-2018)')
plt.xlabel('Year')
plt.ylabel('Packs Sold Per Capita')
plt.legend()
plt.grid(True)
plt.show()

```



```

#Estimate ATEs

#Question 6
#Only 1970-1990, regressing log sales on log price

cig_data = tax_data[(tax_data['Year'] >= 1970) & (tax_data['Year'] <= 1990)]

cig_data['ln_sales'] = np.log(cig_data['sales_per_capita'])
cig_data['ln_price'] = np.log(cig_data['cost_per_pack'])
cig_data['ln_total_tax'] = np.log(cig_data['tax_dollar'])

#running OLS regression
X = sm.add_constant(cig_data['ln_price'])
y = cig_data['ln_sales']

regression_results = sm.OLS(y, X).fit()
print(regression_results.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                  ln_sales    R-squared:                0.126
Model:                            OLS      Adj. R-squared:           0.125
Method:                 Least Squares    F-statistic:                153.9
Date:                 Mon, 17 Mar 2025    Prob (F-statistic):        4.18e-33
Time:                 09:08:01      Log-Likelihood:            148.99
No. Observations:          1071      AIC:                      -294.0
Df Residuals:              1069      BIC:                      -284.0
Df Model:                   1
Covariance Type:            nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|----------|---------|---------|---------|-------|--------|--------|
| const | 4.7504 | 0.008 | 585.321 | 0.000 | 4.734 | 4.766 |
| ln_price | -0.1715 | 0.014 | -12.404 | 0.000 | -0.199 | -0.144 |

```

=====
Omnibus:                 64.611    Durbin-Watson:           0.139
Prob(Omnibus):           0.000    Jarque-Bera (JB):        224.414
Skew:                   0.173    Prob(JB):                1.86e-49
Kurtosis:                5.216    Cond. No.                2.48
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
/var/folders/2q/wzjp_2kd355b8clhzqwmtyb40000gn/T/ipykernel_90894/4233090682.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/1dindexing.html
cig_data['ln_sales'] = np.log(cig_data['sales_per_capita'])
/var/folders/2q/wzjp_2kd355b8clhzqwmtyb40000gn/T/ipykernel_90894/4233090682.py:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/1dindexing.html
cig_data['ln_price'] = np.log(cig_data['cost_per_pack'])
/var/folders/2q/wzjp_2kd355b8clhzqwmtyb40000gn/T/ipykernel_90894/4233090682.py:10: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/1dindexing.html
cig_data['ln_total_tax'] = np.log(cig_data['tax_dollar'])
```

```
#Question 7 and 8, log sales on log prices
#First stage
first_stage = sm.OLS(cig_data['ln_price'], sm.add_constant(cig_data['ln_total_tax'])).fit()
print("First stage Regression (ln_price ~ ln_total_tax):\n")
print(first_stage.summary())
```

First stage Regression (ln_price ~ ln_total_tax):

```

                        OLS Regression Results
=====
Dep. Variable:          ln_price    R-squared:                0.683
Model:                  OLS        Adj. R-squared:             0.683
Method:                 Least Squares    F-statistic:          2301.
Date:                  Mon, 17 Mar 2025    Prob (F-statistic):    8.21e-269
Time:                  09:08:01          Log-Likelihood:        -86.164
No. Observations:      1071             AIC:                  176.3
Df Residuals:          1069             BIC:                  186.3
Df Model:               1
Covariance Type:       nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----

```


| | | | | | | |
|----------------|--------|--------|-------------------|-------|-------|----------|
| const | 1.1786 | 0.033 | 35.712 | 0.000 | 1.114 | 1.243 |
| ln_total_tax | 1.0803 | 0.023 | 47.973 | 0.000 | 1.036 | 1.125 |
| ===== | | | | | | |
| Omnibus: | | 30.760 | Durbin-Watson: | | | 0.408 |
| Prob(Omnibus): | | 0.000 | Jarque-Bera (JB): | | | 32.668 |
| Skew: | | 0.421 | Prob(JB): | | | 8.06e-08 |
| Kurtosis: | | 3.156 | Cond. No. | | | 8.72 |
| ===== | | | | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
# Log prices from the first stage
price_hat = first_stage.predict(sm.add_constant(cig_data['ln_total_tax']))

# Second-stage regression (IV Regression)
second_stage = sm.OLS(cig_data['ln_sales'], sm.add_constant(price_hat)).fit()
print("\nSecond stage Regression (ln_sales ~ pricehat):\n")
print(second_stage.summary())
```

Second stage Regression (ln_sales ~ pricehat):

| OLS Regression Results | | | | | | |
|------------------------|------------------|---------------------|----------------|-------|--------|----------|
| ===== | | | | | | |
| Dep. Variable: | ln_sales | R-squared: | | | | 0.236 |
| Model: | OLS | Adj. R-squared: | | | | 0.235 |
| Method: | Least Squares | F-statistic: | | | | 330.3 |
| Date: | Mon, 17 Mar 2025 | Prob (F-statistic): | | | | 1.56e-64 |
| Time: | 09:08:01 | Log-Likelihood: | | | | 221.17 |
| No. Observations: | 1071 | AIC: | | | | -438.3 |
| Df Residuals: | 1069 | BIC: | | | | -428.4 |
| Df Model: | 1 | | | | | |
| Covariance Type: | nonrobust | | | | | |
| ===== | | | | | | |
| | coef | std err | t | P> t | [0.025 | 0.975] |
| ----- | | | | | | |
| const | 4.7101 | 0.008 | 573.443 | 0.000 | 4.694 | 4.726 |
| 0 | -0.2843 | 0.016 | -18.175 | 0.000 | -0.315 | -0.254 |
| ===== | | | | | | |
| Omnibus: | | 83.338 | Durbin-Watson: | | | 0.157 |

| | | | |
|----------------|-------|-------------------|----------|
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 430.014 |
| Skew: | 0.023 | Prob(JB): | 4.20e-94 |
| Kurtosis: | 6.104 | Cond. No. | 2.98 |

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
#Question 9 - repeat for 1991-2015
cig_data2 = tax_data[(tax_data['Year'] >= 1991) & (tax_data['Year'] <= 2015)]

cig_data2['ln_sales'] = np.log(cig_data2['sales_per_capita'])
cig_data2['ln_price'] = np.log(cig_data2['cost_per_pack'])
cig_data2['ln_total_tax'] = np.log(cig_data2['tax_dollar'])

#running OLS regression
X2 = sm.add_constant(cig_data2['ln_price'])
Y2 = cig_data2['ln_sales']

reg2 = sm.OLS(Y2, X2).fit()
print(reg2.summary())
```

OLS Regression Results

=====

| | | | |
|-------------------|------------------|---------------------|-----------|
| Dep. Variable: | ln_sales | R-squared: | 0.533 |
| Model: | OLS | Adj. R-squared: | 0.532 |
| Method: | Least Squares | F-statistic: | 1451. |
| Date: | Mon, 17 Mar 2025 | Prob (F-statistic): | 1.52e-212 |
| Time: | 09:08:01 | Log-Likelihood: | -296.47 |
| No. Observations: | 1275 | AIC: | 596.9 |
| Df Residuals: | 1273 | BIC: | 607.2 |
| Df Model: | 1 | | |
| Covariance Type: | nonrobust | | |

=====

| | coef | std err | t | P> t | [0.025 | 0.975] |
|----------|---------|---------|---------|-------|--------|--------|
| const | 5.0395 | 0.023 | 219.934 | 0.000 | 4.995 | 5.084 |
| ln_price | -0.6656 | 0.017 | -38.094 | 0.000 | -0.700 | -0.631 |

=====

| | | | |
|----------------|--------|-------------------|--------|
| Omnibus: | 19.351 | Durbin-Watson: | 0.158 |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 33.046 |

| | | | |
|-----------|-------|-----------|----------|
| Skew: | 0.064 | Prob(JB): | 6.67e-08 |
| Kurtosis: | 3.778 | Cond. No. | 5.37 |

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
/var/folders/2q/wzjp_2kd355b8clhzqwytb40000gn/T/ipykernel_90894/1083457502.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/1d.html#copying
cig_data2['ln_sales'] = np.log(cig_data2['sales_per_capita'])
/var/folders/2q/wzjp_2kd355b8clhzqwytb40000gn/T/ipykernel_90894/1083457502.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/1d.html#copying
cig_data2['ln_price'] = np.log(cig_data2['cost_per_pack'])
/var/folders/2q/wzjp_2kd355b8clhzqwytb40000gn/T/ipykernel_90894/1083457502.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/1d.html#copying
cig_data2['ln_total_tax'] = np.log(cig_data2['tax_dollar'])
```

```
#QUESTION 10
# Creating a table to summarize results
summary_table = pd.DataFrame({
    'Years': ['1970-1990', '1970-1990', '1991-2015', '1991-2015'],
```

ValueError: All arrays must be of the same length

ValueError Traceback (most recent call last)

Input In [29], in <cell line: 3>()

```
1 #QUESTION 10
2 # Creating a table to summarize results
----> 3 results_table = pd.DataFrame(      4      'Time Period': ['1970-1990', '1970-1990', '1991-2015', '1991-2015'],
13 # Print the table without the index and with lines between each column/row
14 print(results_table.to_string(index=False, line_width=80))
File ~/anaconda/lib/python3.9/site-packages/pandas/core/frame.py:636, in DataFrame.__init__(self, data, index, columns, dtype, copy)
```

```

630     mgr = self._init_mgr(
631         data, axes="index": index, "columns": columns, dtype=dtype, copy=copy
632     )
634 elif isinstance(data, dict):
635     # GH#38939 de facto copy defaults to False only in non-dict cases
--> 636     mgr = dict_to_mgr(data, index, columns, dtype=dtype, copy=copy, typ=manager)
637 elif isinstance(data, ma.MaskedArray):
638     import numpy.ma.mrecords as mrecords
File ~/anaconda/lib/python3.9/site-packages/pandas/core/internals/construction.py:502, in di
494     arrays = [
495         x
496         if not hasattr(x, "dtype") or not isinstance(x.dtype, ExtensionDtype)
497         else x.copy()
498         for x in arrays
499     ]
500     # TODO: can we get rid of the dt64tz special case above?
--> 502 return arrays_to_mgr(arrays, columns, index, dtype=dtype, typ=typ, consolidate=copy)
File ~/anaconda/lib/python3.9/site-packages/pandas/core/internals/construction.py:120, in ar
117 if verify_integrity:
118     # figure out the index, if necessary
119     if index is None:
--> 120         index = _extract_index(arrays)
121     else:
122         index = ensure_index(index)
File ~/anaconda/lib/python3.9/site-packages/pandas/core/internals/construction.py:674, in _e
672 lengths = list(set(raw_lengths))
673 if len(lengths) > 1:
--> 674     raise ValueError("All arrays must be of the same length")
676 if have_dicts:
677     raise ValueError(
678         "Mixing dicts with non-Series may lead to ambiguous ordering."
679     )
ValueError: All arrays must be of the same length

```